



Asia's Largest

Cloud & AI

Conference 2023

17 - 18, November 2023
IIT Madras Research Park, Chennai



Amazon DynamoDB Deep Dive

Architecting Data Models for High-Performance Applications

4 people walk into a NOSQL restaurant...
but immediately leaves why?

Because they couldn't find a table!





```
{  
  name: "Bhuvana",  
  role: "Technical Architect",  
  company: "Accenture",  
  skills: [ "NodeJS", "AWS", "SQL", "NoSQL"],  
  volunteer: ["AWS Community Builder", "AZConf 2023"],  
  sideHustles: ["Tech Speaker", "Runs a Children's Library"],  
  favoriteMovies: ["Kungfu Panda", "Soul"],  
  favoriteQuote: "Just Keep Swimming!"  
}
```

Agenda

What is DynamoDB?

Why DynamoDB?

How DynamoDB?

"Computers are useless. They can only give you answers."

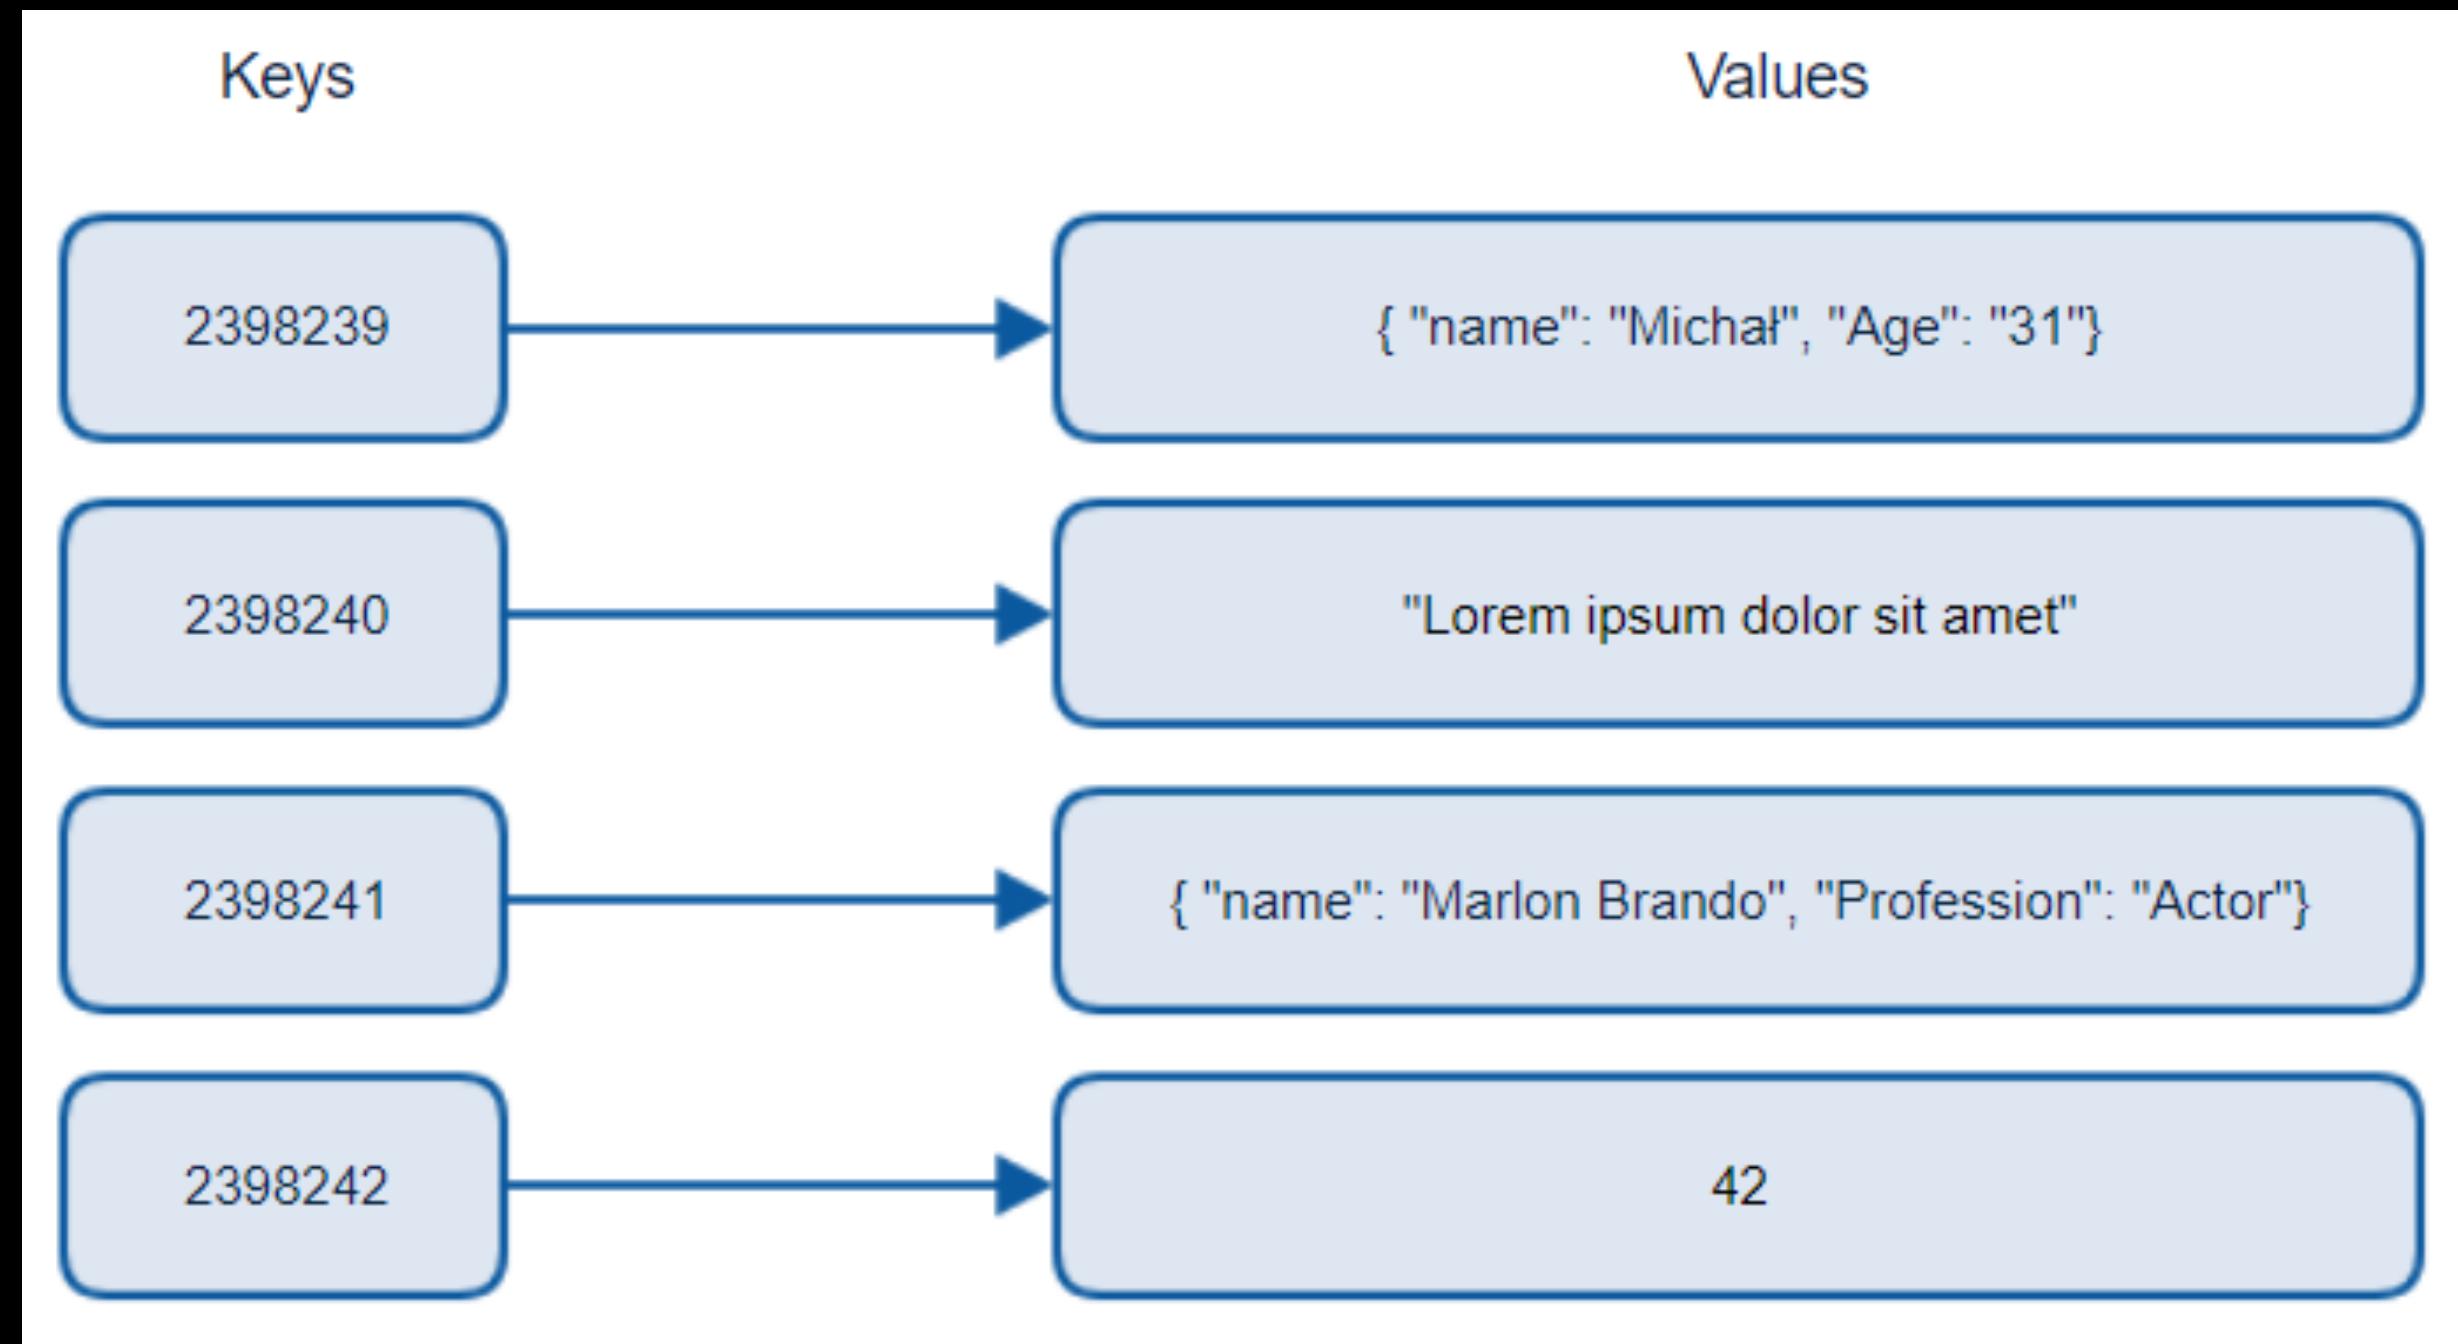
– Pablo Picasso, 1968

What is DynamoDB

- Amazon DynamoDB is a fully managed NoSQL database service that provides fast and predictable performance with seamless scalability. DynamoDB lets you offload the administrative burdens of operating and scaling a distributed database so that you don't have to worry about hardware provisioning, setup and configuration, replication, software patching, or cluster scaling.
- With DynamoDB, you can create database tables that can store and retrieve any amount of data and serve any level of request traffic. You can scale up or scale down your tables' throughput capacity without downtime or performance degradation. You can use the AWS Management Console to monitor resource utilization and performance metrics.

NoSQL Key-Value or Document Store

What?



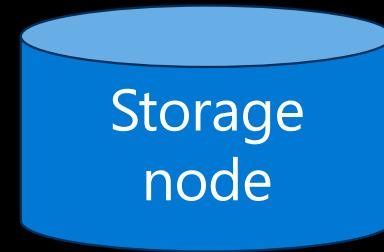


Predictable, Single digit millisecond response time at any scale

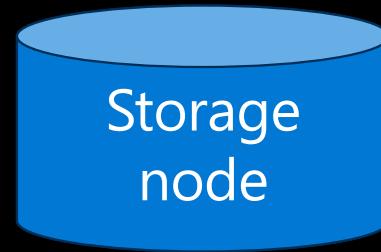
How?

Partition Key

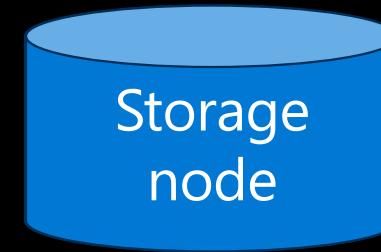
userId	created_ts	orderId	orderTotal
user1	01/07/23	order#1	2000
user2	03/07/23	order#2	600
user3	05/07/23	order#3	100
user4	01/07/23	order#4	1500
user5	05/07/23	order#5	1242



user#1



user#2
user#3



user#4
user#5



Horizontal Scaling - Partition size limited to about 10GB

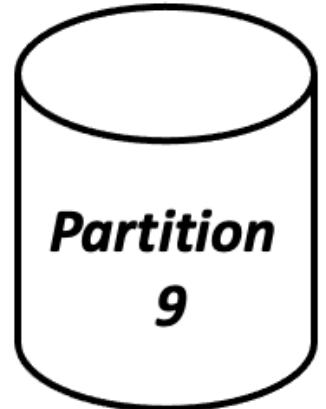
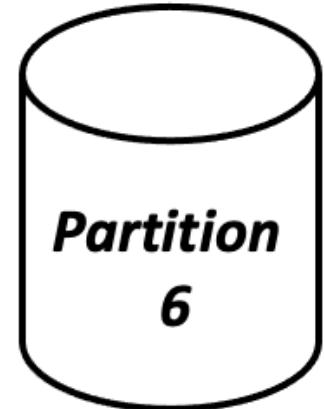
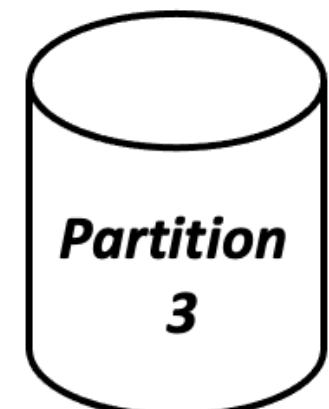
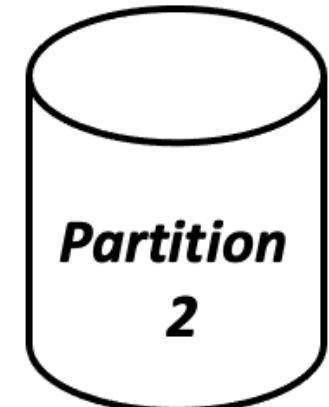
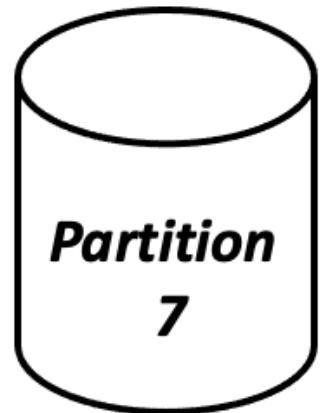
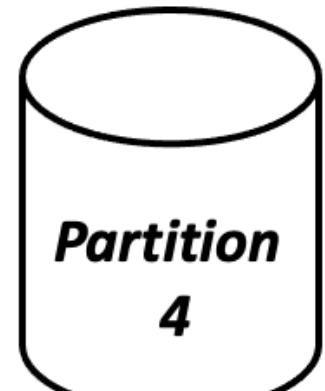
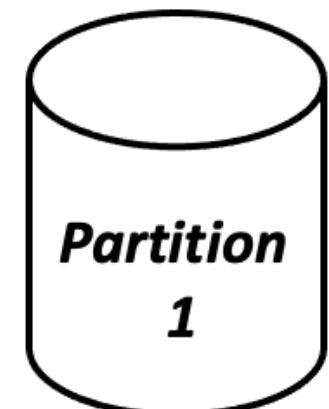
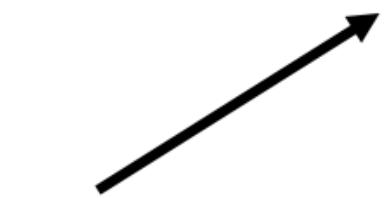
CustomerId: "de91538a"

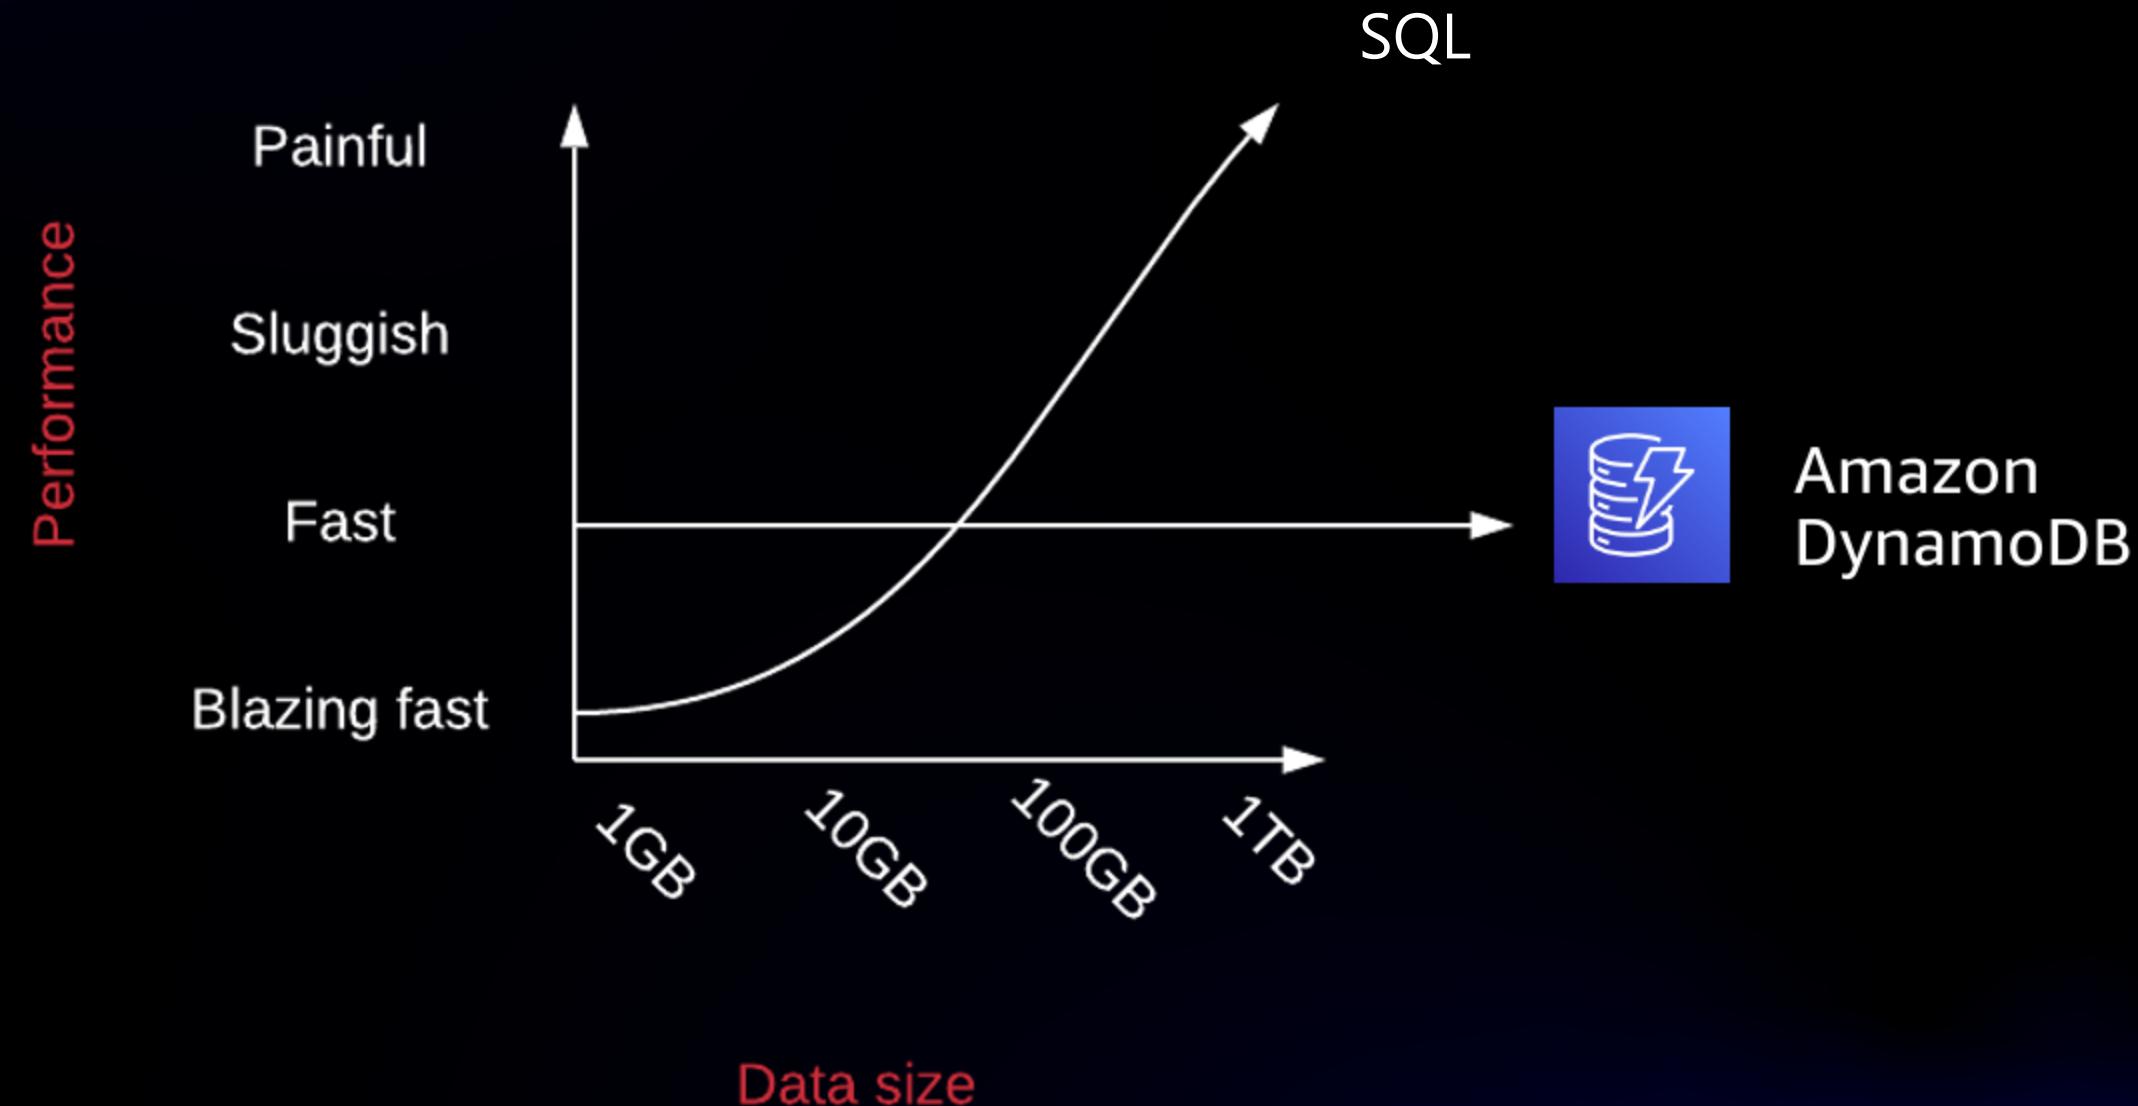


Request router

$O(1)$

*fx(CustomerId):
This item belongs to
Partition 1*







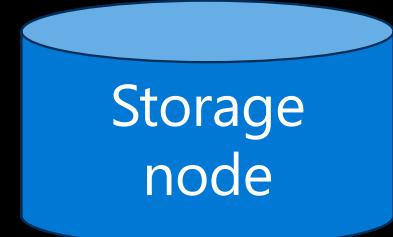
Highly Available - 99.999% Availability

How?

Availability Zone 1



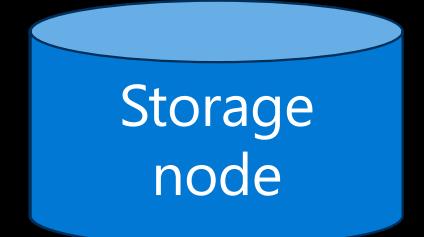
user#1



Storage
node

user#2

user#3

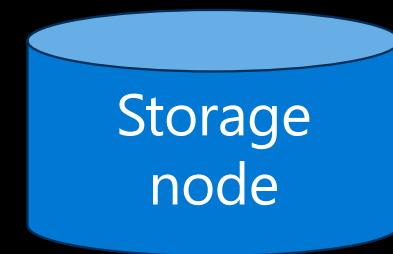


Storage
node

user#4

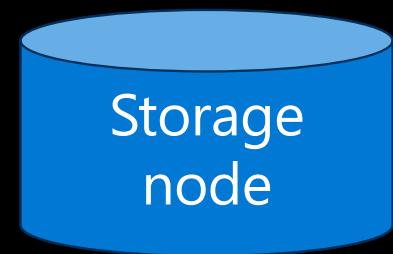
user#5

Availability Zone 2



Storage
node

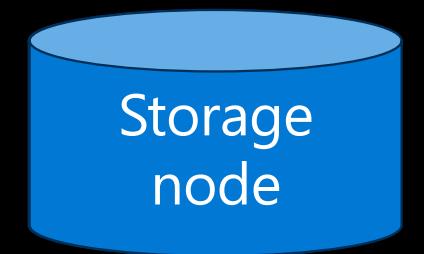
user#1



Storage
node

user#2

user#3

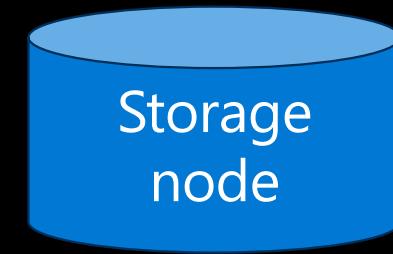


Storage
node

user#4

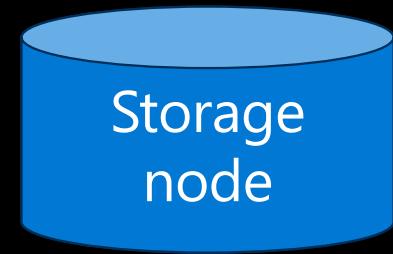
user#5

Availability Zone 3



Storage
node

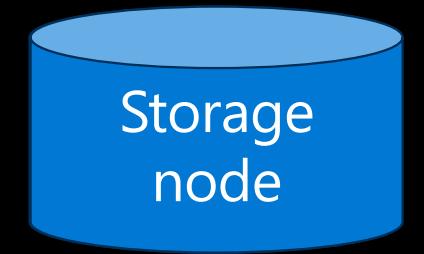
user#1



Storage
node

user#2

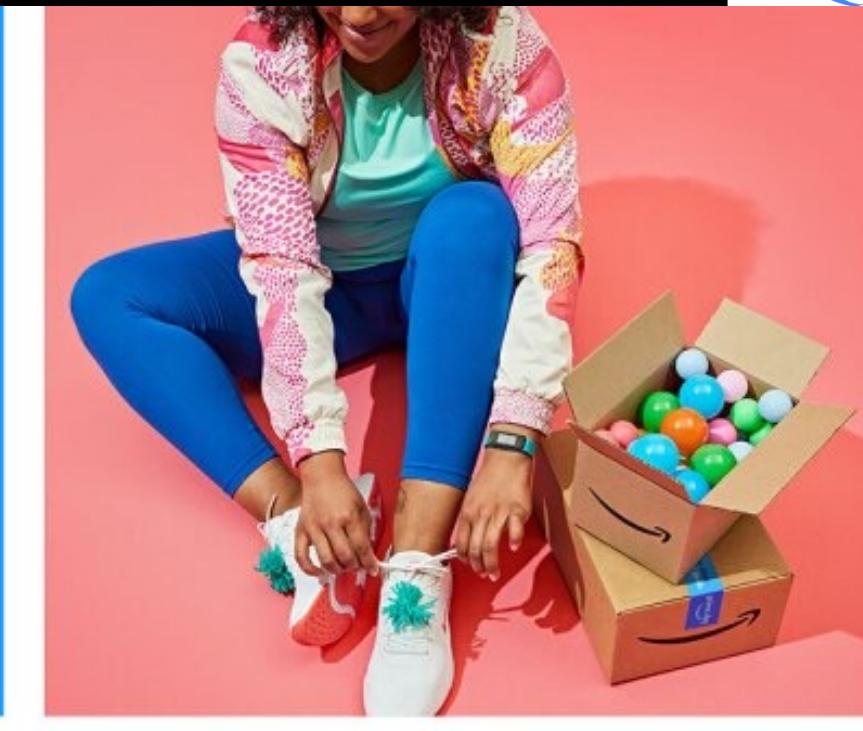
user#3



Storage
node

user#4

user#5



Amazon DynamoDB – DynamoDB powers multiple high-traffic Amazon properties and systems including [Alexa](#), the [Amazon.com](#) sites, and all [Amazon fulfillment centers](#). Over the course of Prime Day, these sources made trillions of calls to the DynamoDB API. DynamoDB maintained high availability while delivering single-digit millisecond responses and peaking at 126 million requests per second.



Serverless

What?

Serverless

- A fully managed service –
No worry about things like server health, storage, and network connectivity
- Infinitely scalable





Pricing?

Pricing

- Pay-per-use model
 - you never pay for hardware or services you're not actually using
- Traditional DB – CPU, RAM, IOPS
- Pay for reads and writes – RCU and WCU

Why are we discussing architecture?

To model data that exploits this architecture

Why we need to understand architecture?

To get consistent, predictable and reliable performance from DynamoDB

Terminology

- Table
- Items – (SQL Rows)
- Primary Key
- Attributes – (SQL Columns)

Table - Employee

Primary key				
Partition key: LoginAlias				
Attributes				
johns	firstName	LastName	ManagerLoginAlias	Skills
	John	Stiles	NA	["executive management"]
marthar	firstName	LastName	ManagerLoginAlias	Skills
	Martha	Rivera	johns	["software", "management"]
mateoj	firstName	LastName	ManagerLoginAlias	Skills
	Mateo	Jackson	marthar	["software"]
janed	firstName	LastName	ManagerLoginAlias	Skills
	Jane	Doe	marthar	["software"]
diegor	firstName	LastName	ManagerLoginAlias	Skills
	Diego	Ramirez	johns	["executive assistant"]

Primary Key

- Simple Primary Key (Partition Key)
- Composite Primary Key (Partition Key + Sort Key)

How to choose Primary Key?

- Partition key is for grouping
- Sort key is for ordering

Partition Key

- High cardinality to distribute requests



PutItem:

CustomerId: "de91538a"

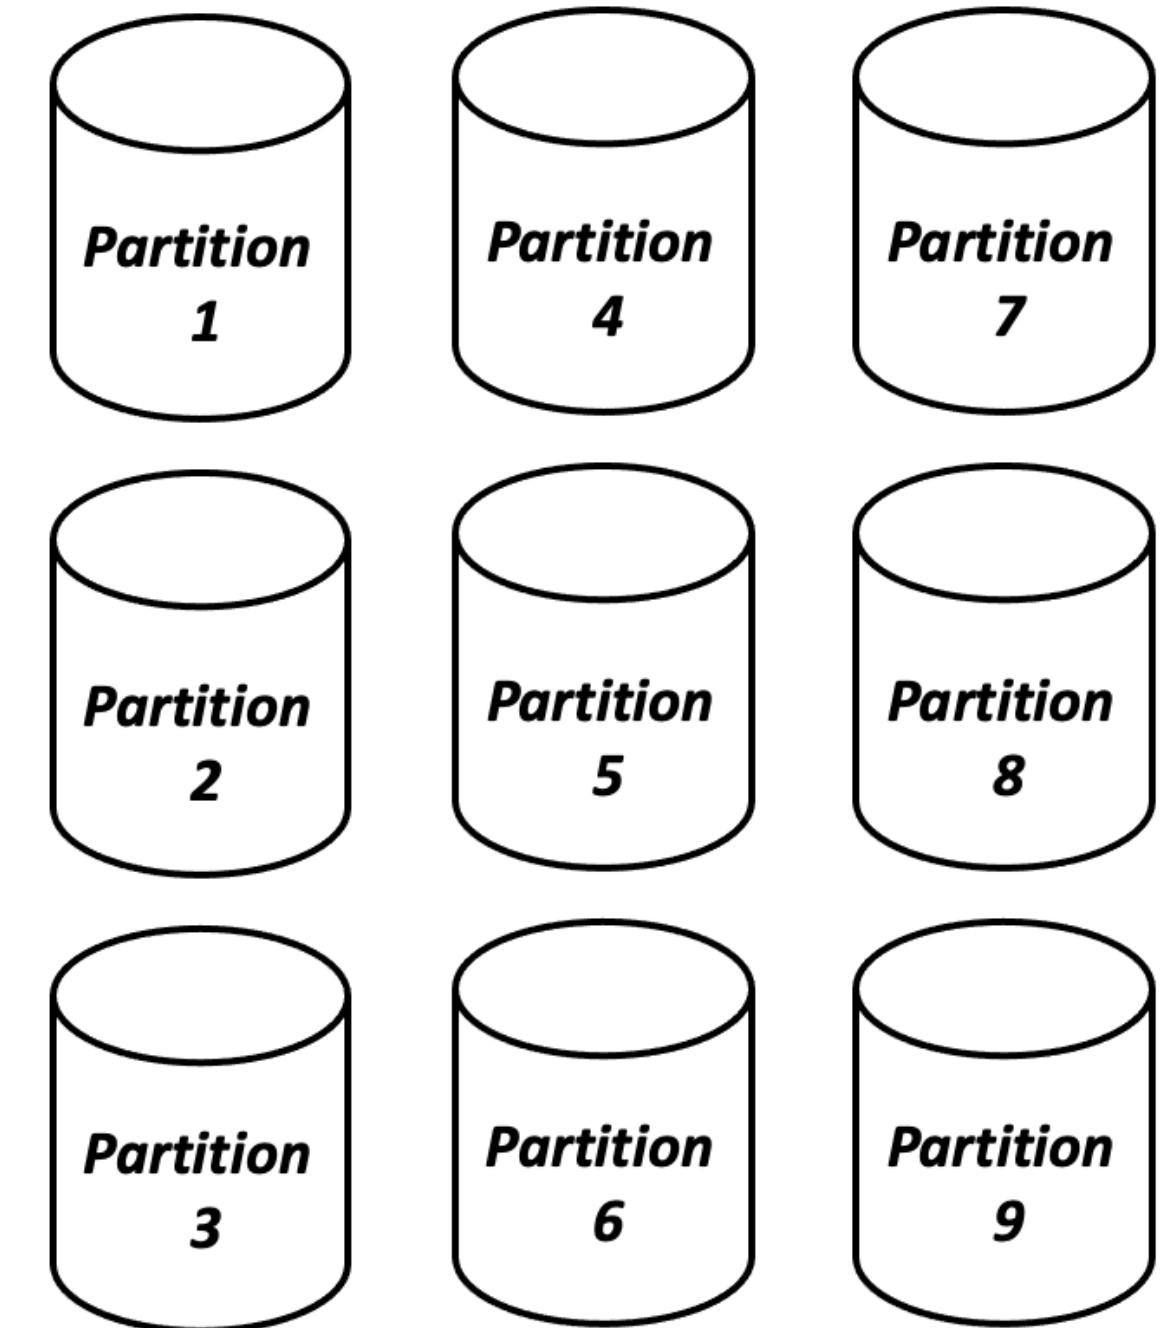
OrderId: "298b628e"



Request router

$O(1)$

fx(CustomerId):
This item belongs to
Partition 1





PutItem:

CustomerId: "de91538a"

OrderId: "298b628e"

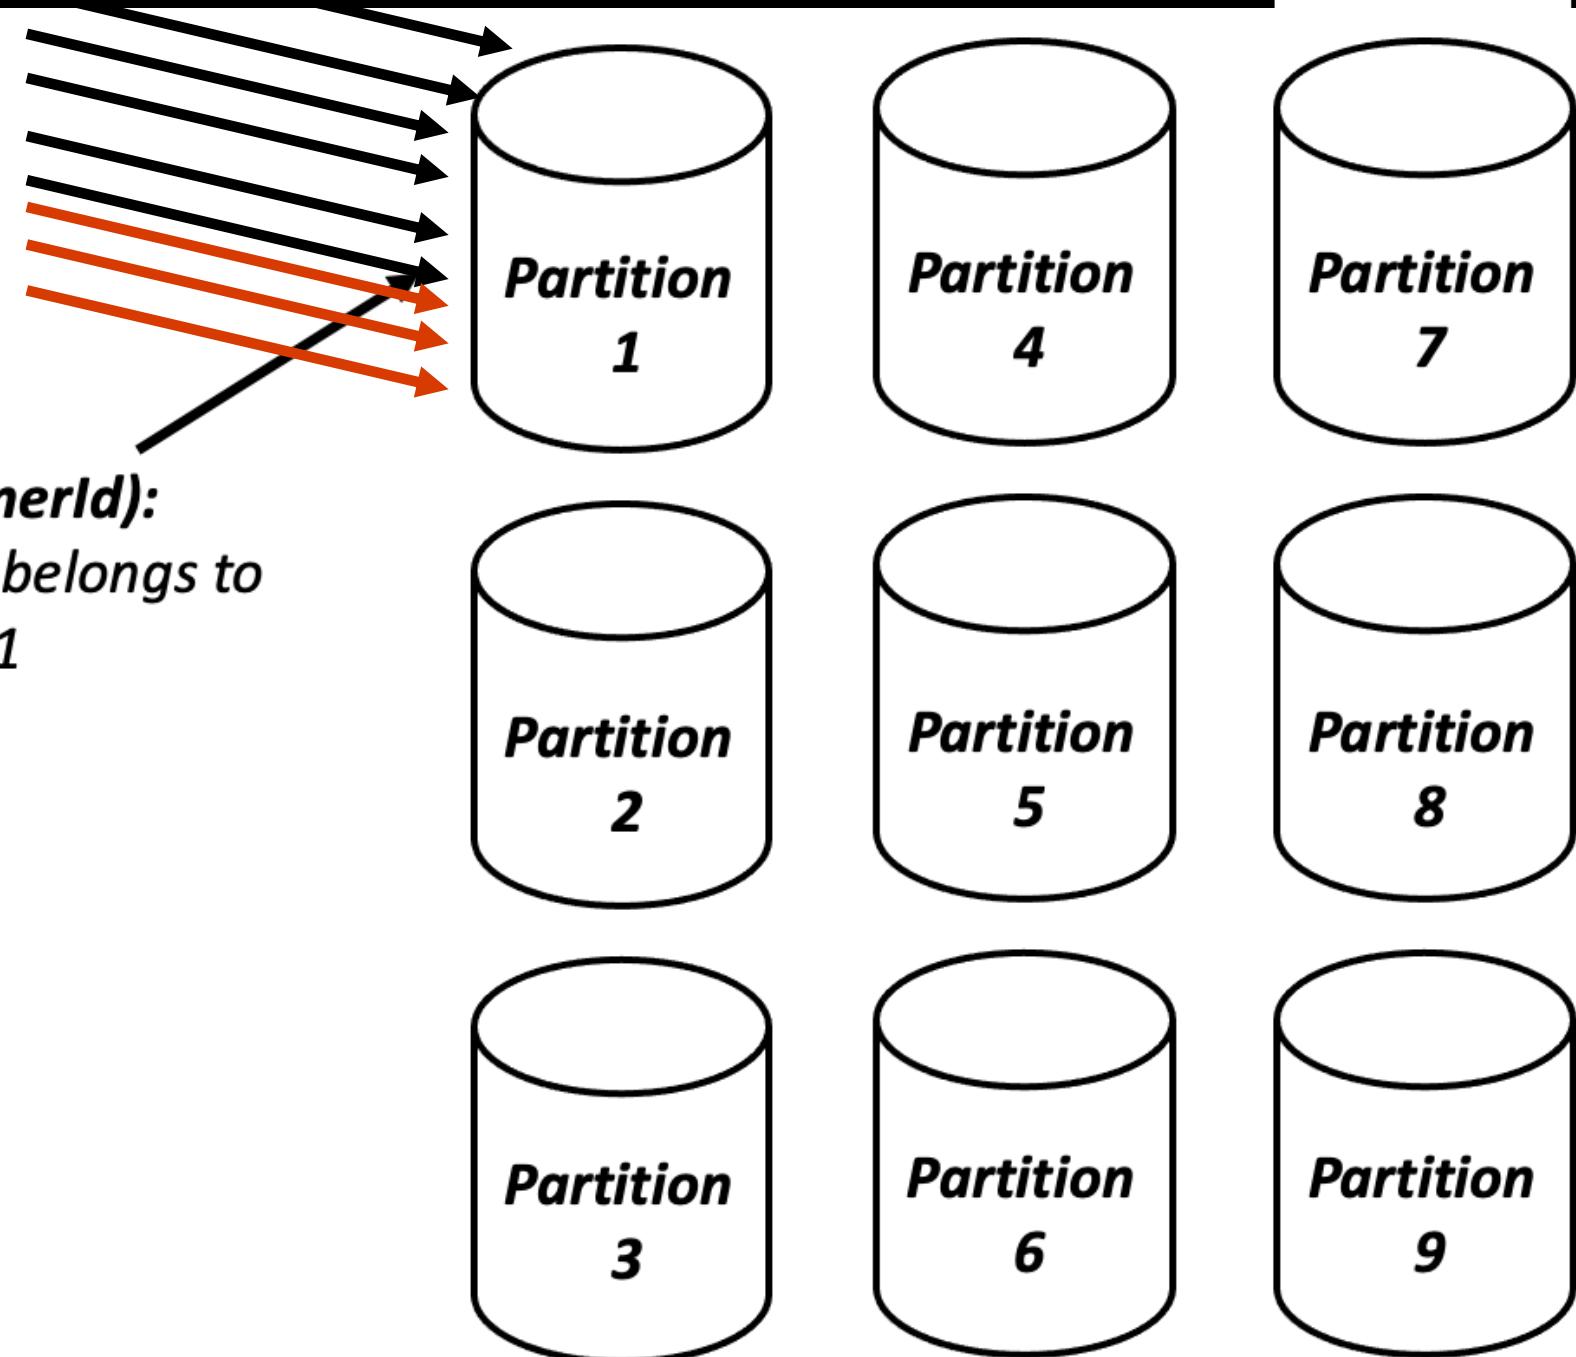


Request router

$O(1)$

3,000 RCU/s
1,000 WCU/s

fx(CustomerId):
This item belongs to
Partition 1



Partition Key

- High cardinality to distribute requests
- Unique and meaningful
- Cannot update the primary key – field nor value

Sort Key

- For Ordering
- Alphabetical
- ISO Timestamps
- UUID - Universally Unique Lexicographically Sortable Identifier

Sort Key Operations

- $a = b$
- $a < b$
- $a \leq b$
- $a > b$
- $a \geq b$
- a BETWEEN b AND c
- begins_with (a, substr)

Choosing Primary Key is one of the most important decisions in the DynamoDB data modeling.



DynamoDB Read and Write Operations

Read Operations

Query	Scan
Retrieves items based on specified key criteria	Scans the entire table or a subset based on filters
Highly efficient for retrieving specific items	Less efficient for large tables or full-table scans
Can query based on partition key and sort key	Limited to filtering based on non-key attributes
Consumes fewer read capacity units (RCUs)	Consumes more RCUs due to scanning the entire table
Lower cost as it retrieves specific items	Higher cost due to scanning the entire or large subset
Ideal for retrieving specific items by key	Suitable for ad-hoc queries or infrequent use cases

Write Operations

- PutItem
- UpdateItem
- DeleteItem
- BatchWriteItem

**What if we want to query based on attributes
that are not primary key?**



Indexing

Local Secondary Index

TABLE	Id (Partition)	Name (Sort Key)	Address	Phone	Email		
LSI	Id (Partition)	Address (Sort Key)	Name (Key)			KEYS ONLY	
	Id (Partition)	Phone (Sort Key)	Name (Key)	Email		INCLUDE "EMAIL"	Projections
	Id (Partition)	Email (Sort Key)	Name (Key)	Address	Phone	ALL	

Global secondary index

TABLE	Id (Partition)	Name	Address	Phone	Email		
GSI	Name (Partition)	Id (Key)				KEYS ONLY	Projections
	Phone (Partition)	Address (Sort Key)	Id (Key)	Name		INCLUDE "NAME"	
	Email (Partition)	Phone (Sort Key)	Id (Key)	Name	Address	ALL	

Demo





Data Model Blog App

CLOUD INFORMATION



Lorum ipsum Follow 64

Lorum ipsum dolor sit amet, aliqua.
Duis aute irure dolor in reprehenderit in voluptate
velit esse cillum dolore eu fugiat nulla pariatur.



Lorum ipsum Follow 64

Lorum ipsum dolor sit amet, aliqua.
Duis aute irure dolor in reprehenderit in voluptate
velit esse cillum dolore eu fugiat nulla pariatur.



Lorum ipsum Follow 64

Lorum ipsum dolor sit amet, aliqua.
Duis aute irure dolor in reprehenderit in voluptate
velit esse cillum dolore eu fugiat nulla pariatur.



Lorum ipsum Follow 64

Lorum ipsum dolor sit amet, aliqua.
Duis aute irure dolor in reprehenderit in voluptate
velit esse cillum dolore eu fugiat nulla pariatur.



Lorum ipsum Follow 64

Lorum ipsum dolor sit amet, aliqua.
Duis aute irure dolor in reprehenderit in voluptate
velit esse cillum dolore eu fugiat nulla pariatur.



1:59

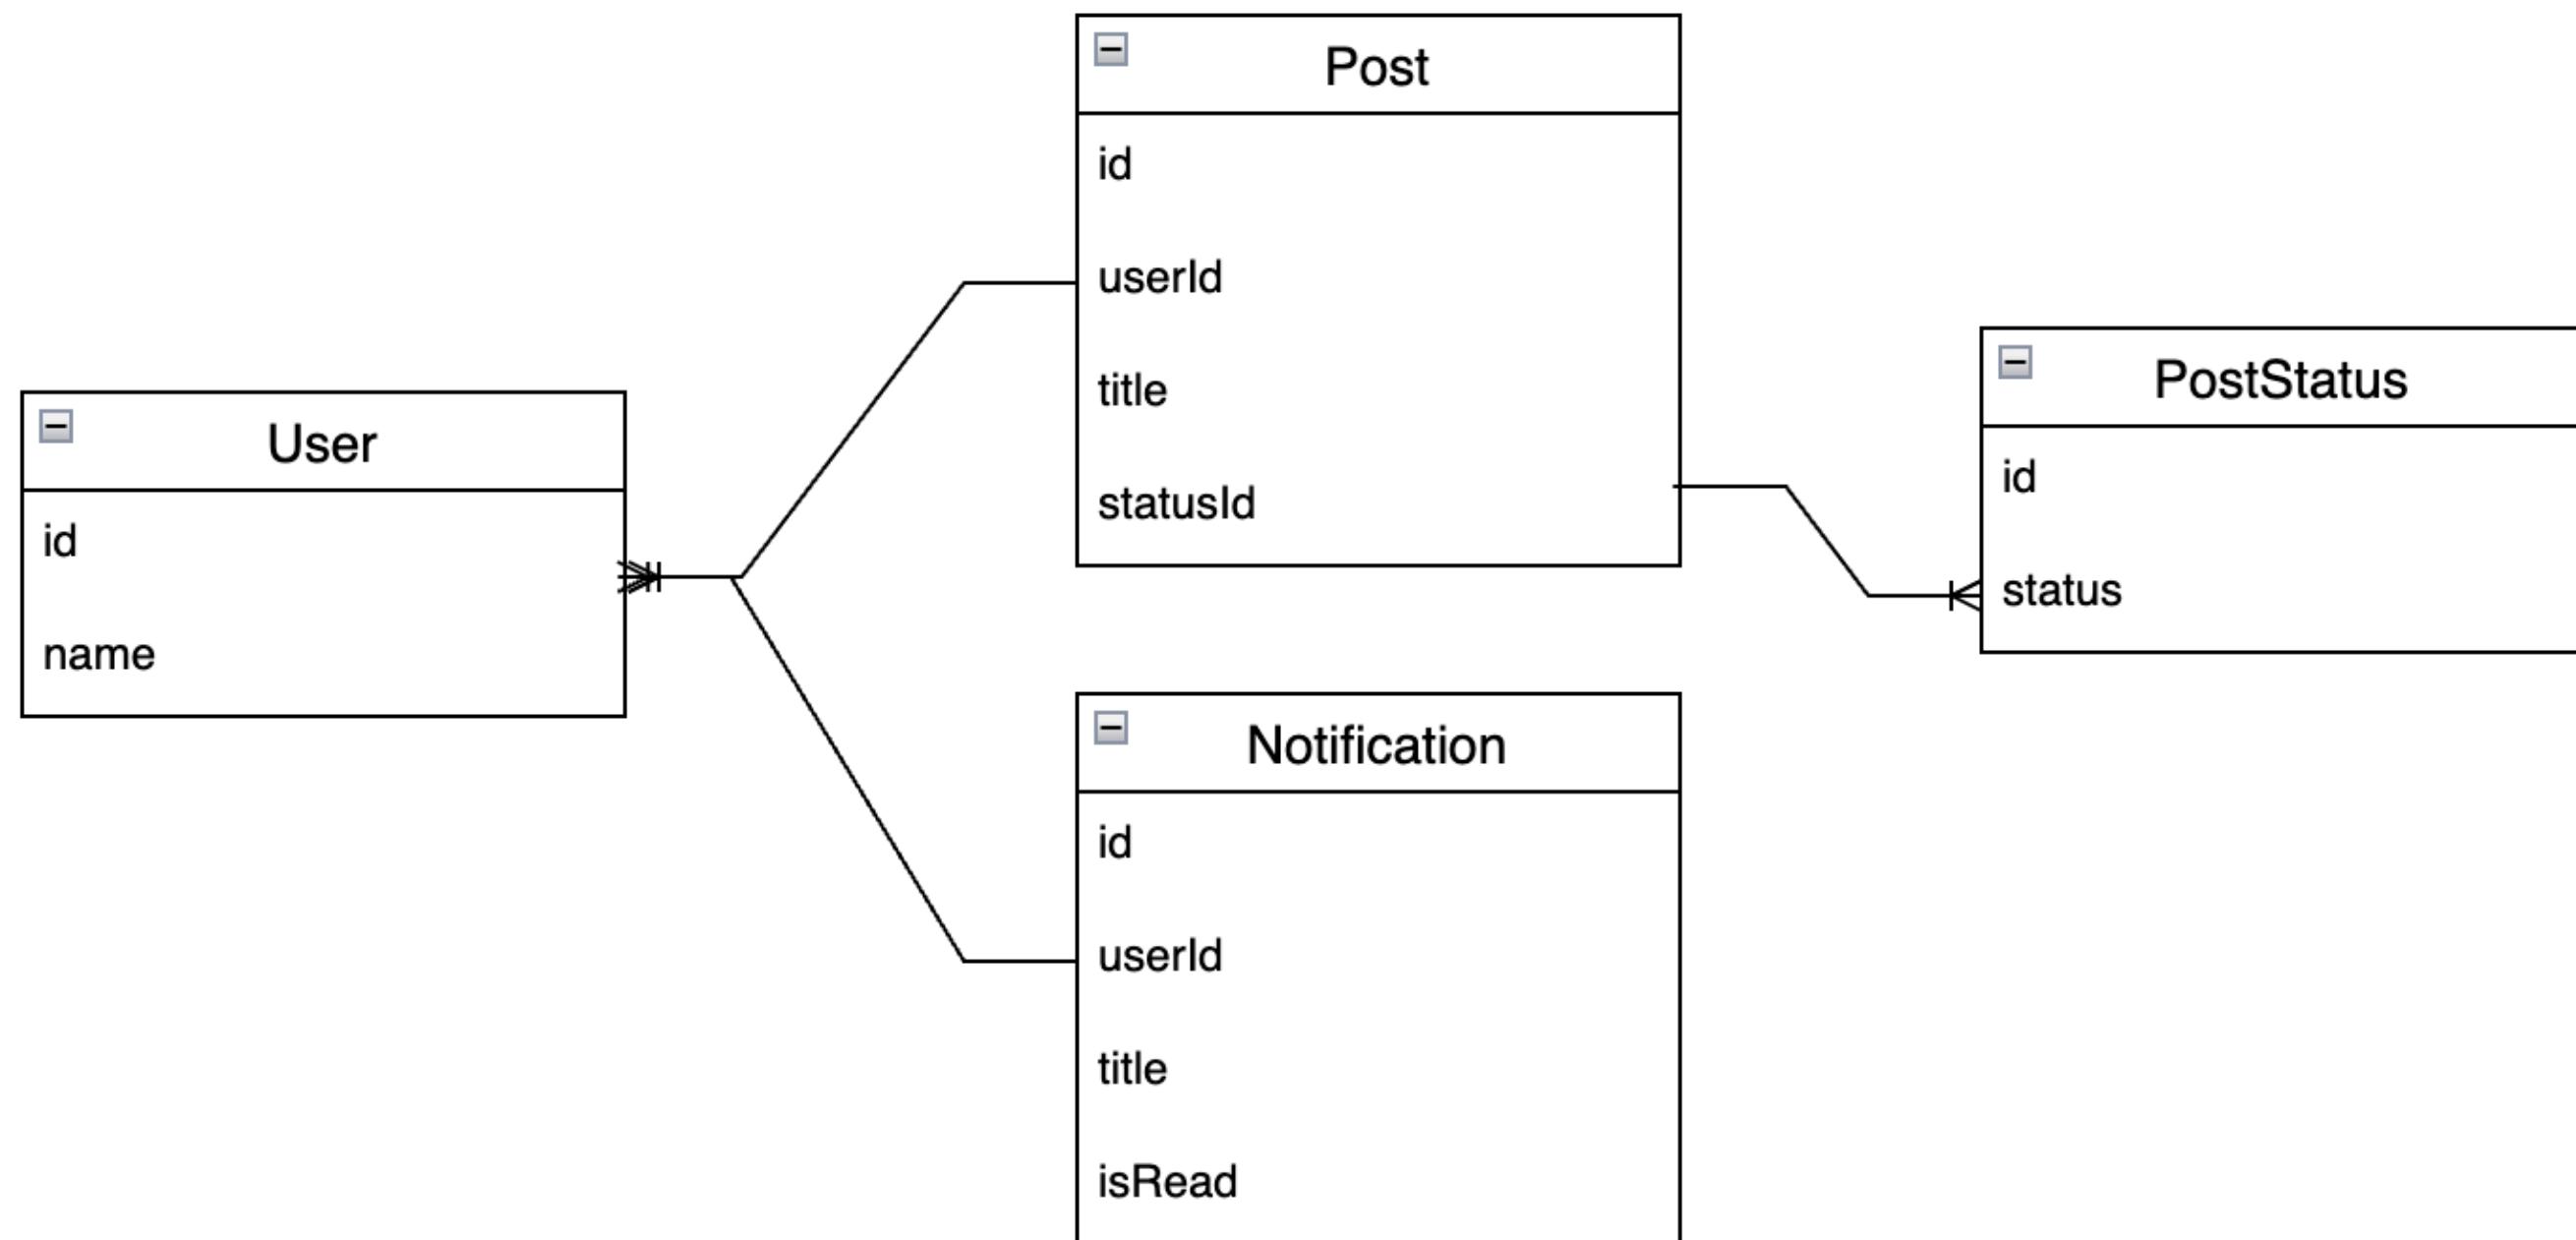


-   Emil Krajewski followed your work
-  Honza Růžička followed your work
-  Fire Media followed your work
-  Alterplay Ltd followed your work
-  Tomas Hrabec followed your work
-  Fappi gen followed your work
-  Christina Sidhoum followed your work
-  Jason Zhang followed your work
-  Jason Zhang followed your work





Entity Relationship Diagram





Schema for DynamoDB

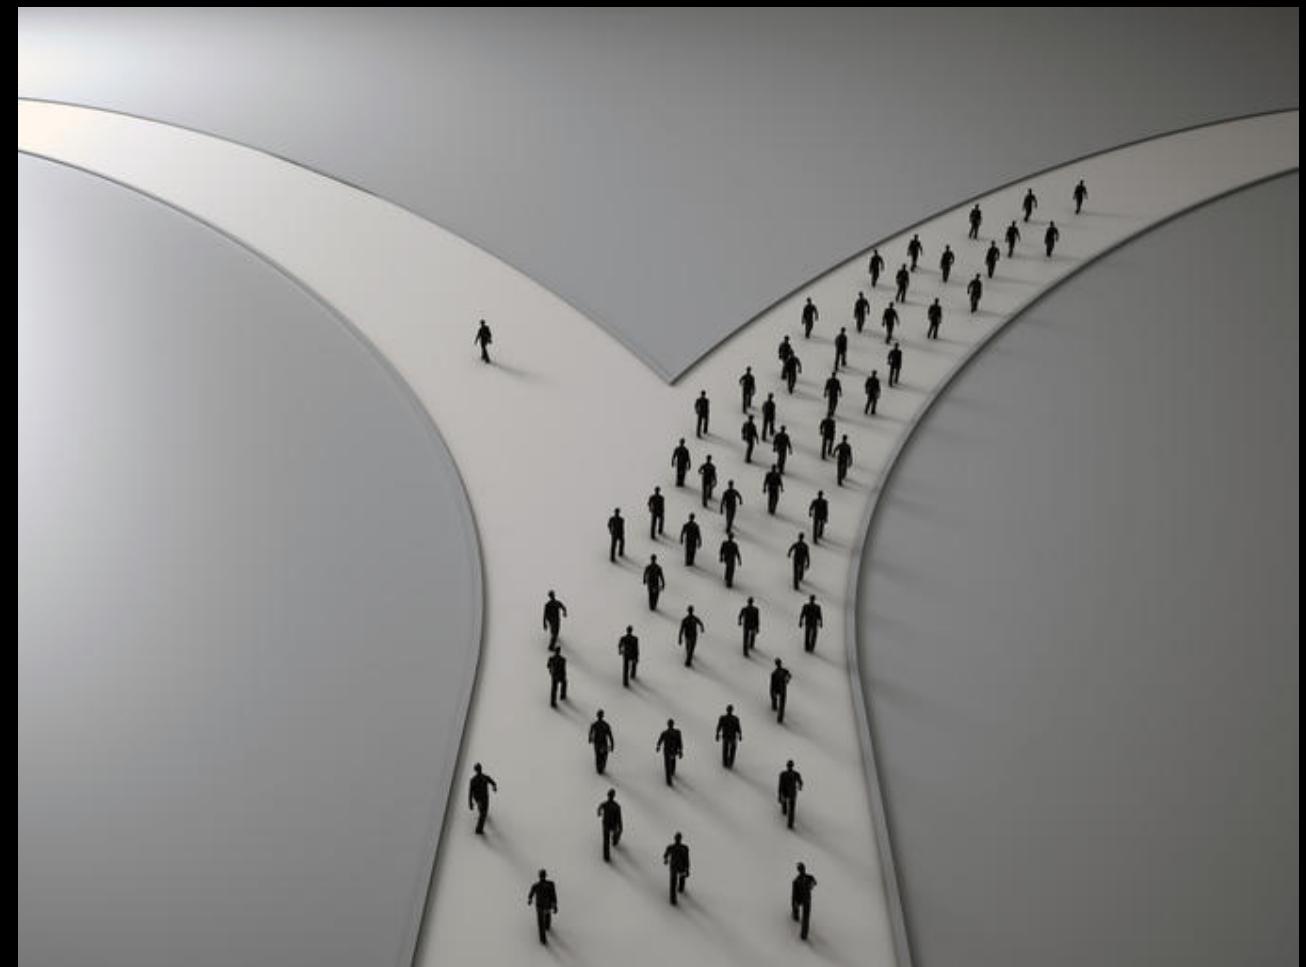
Relational
Data
Schema

NoSQL
Data
Schema

Corporate needs you to find the differences
between this picture and this picture.

They're the same picture.

Different Mindset

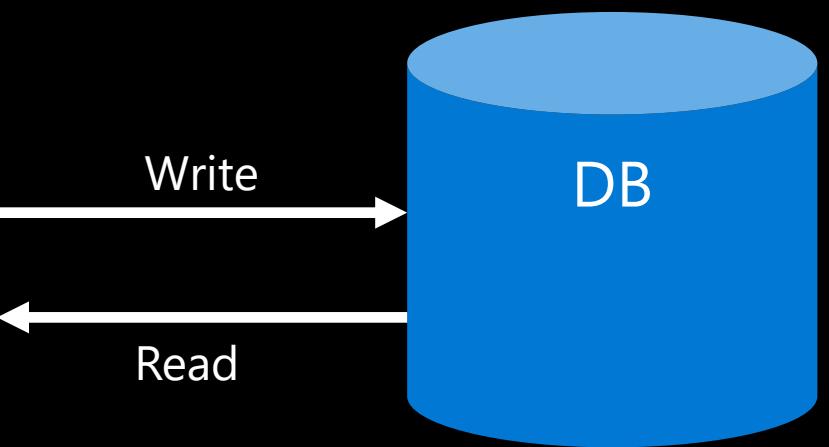


How NoSQL design is different

- By contrast, you shouldn't start designing your schema for DynamoDB until you know the questions it will need to answer. Understanding the business problems and the application use cases up front is essential.
- You should maintain as few tables as possible in a DynamoDB application. Having fewer tables keeps things more scalable, requires less permissions management, and reduces overhead for your DynamoDB application. It can also help keep backup costs lower overall.

Access Patterns

How you access your data – Business Usecases



Write Operations

1. Create User
2. Create Post
3. Update Post
4. Like a Post
5. Create Notification
6. Mark Notification As Read

Read Operations

1. Get User profile
2. Get All Posts for a user
3. Get All Published Posts for a user
4. Get Draft Posts for a user
5. Sort Published Posts By Likes and return top 3 for featured section
6. Get No. of Published Posts for a user
7. Get All Notifications for a user
8. Get Unread Notifications for a user
9. Get No. of Unread Notifications for a user



**WHY DO NOSQL
DEVELOPERS
EAT LUNCH ALONE?**



I DON'T KNOW, WHY? THEY DON'T KNOW HOW TO JOIN TABLES

Single Table Design

WHAT!!





1. Get User profile

Primary key		Attributes	
		Name	
Partition key: userId		Sort key: sk	
user_1	profile	Name	
user_2	profile	Name	
user_3	profile	Name	
		Rose Luettgen	
		Ida Hegmann Sr.	
		Ralph Walker	

2. Get All Posts for a user

{userId: "user1", sk: {BEGINS_WITH : "post#" } }

Primary key		Attributes		
Partition key: userId	Sort key: sk	Title	Status	NoOfLikes
user_1	post#1	Title of Post 1	PUBLISHED	10
		Title of Post 2	DRAFT	
	post#3	Title of Post 3	PUBLISHED	7
		Title of Post 4	PUBLISHED	12
	profile	Name	Rose Luetten	

3. Get All Published Posts for a user

Filter: { Status: PUBLISHED }

Primary key		Attributes		
Partition key: userId	Sort key: sk	Title	Status	NoOfLikes
user_1	post#1	Title	Status	NoOfLikes
		Title of Post 1	PUBLISHED	10
	post#2	Title	Status	NoOfLikes
		Title of Post 2	DRAFT	10
	post#3	Title	Status	NoOfLikes
		Title of Post 3	PUBLISHED	7
	post#4	Title	Status	NoOfLikes
		Title of Post 4	PUBLISHED	12
	profile	Name		
		Rose Luettgen		



4. Get All Draft Posts for a user

Filter: { Status: DRAFT }

Primary key		Attributes		
Partition key: userId	Sort key: sk	Title	Status	NoOfLikes
user_1	post#1	Title	Status	NoOfLikes
		Title of Post 1	PUBLISHED	10
	post#2	Title	Status	
		Title of Post 2	DRAFT	
	post#3	Title	Status	NoOfLikes
		Title of Post 3	PUBLISHED	7
	post#4	Title	Status	NoOfLikes
		Title of Post 4	PUBLISHED	12
	profile	Name		
		Rose Luettgen		

Index - { userId , status }

5. Sort Published Posts By Likes and return top 3 for featured section

Index - { userId , noOfLikes }

Solution

1. Add NoOfLikes = -1 for Draft Posts

2. Create Index on { userId, NoOfLikes }

Primary key		Attributes		
Partition key: userId	Sort key: sk			
user_1	post#1	Title	Status	NoOfLikes
		Title of Post 1	PUBLISHED	10
	post#2	Title	Status	NoOfLikes
		Title of Post 2	DRAFT	-1
	post#3	Title	Status	NoOfLikes
		Title of Post 3	PUBLISHED	7
	post#4	Title	Status	NoOfLikes
		Title of Post 4	PUBLISHED	12
user_2	profile	Name	NoOfPublishedPosts	
		Rose Luettgen	3	
	post#10	Title	Status	NoOfLikes
		Title of Post 10	PUBLISHED	0
	post#11	Title	Status	NoOfLikes
		Title of Post 11	DRAFT	-1
	profile	Name	NoOfPublishedPosts	
		Ida Hegmann Sr.	1	

noOfLikesIndex

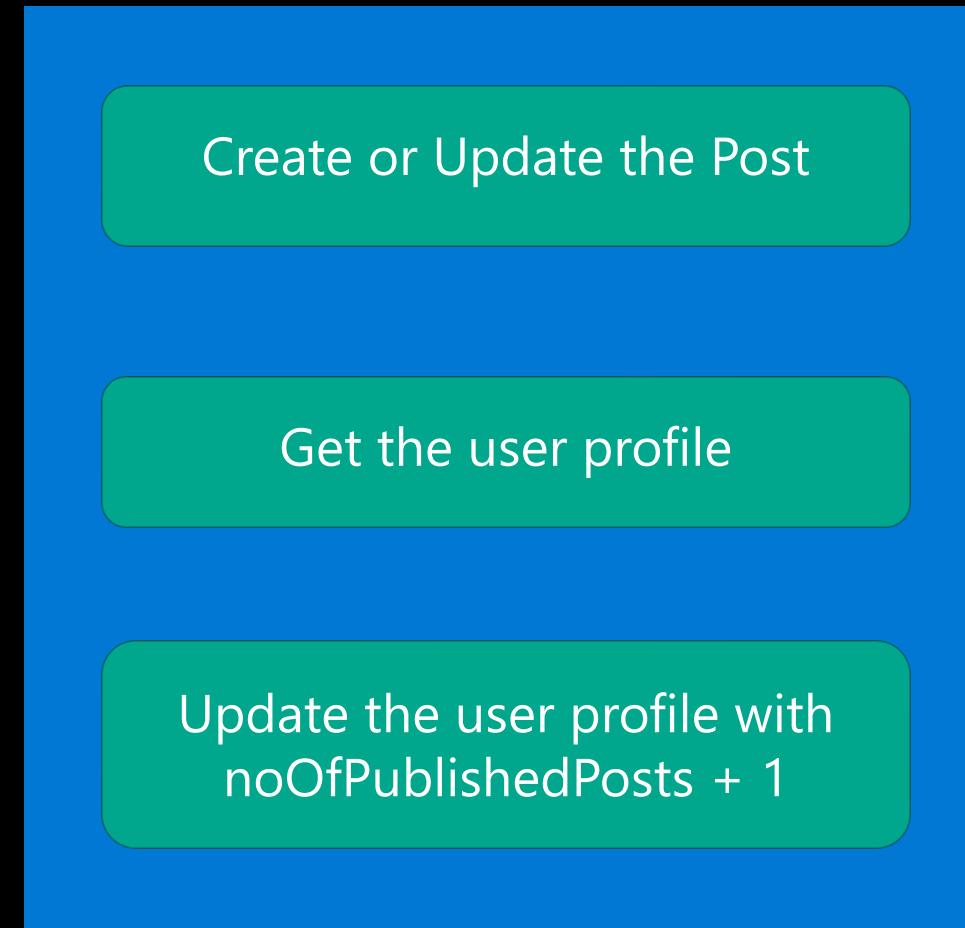
Primary key		Attributes			
	Partition key: userId	Sort key: NoOfLikes	sk	Title	Status
user_1		-1	post#2	Title of Post 2	DRAFT
		10	post#1	Title of Post 1	PUBLISHED
		12	post#4	Title of Post 4	PUBLISHED
		7	post#3	Title of Post 3	PUBLISHED
user_2		-1	post#11	Title of Post 11	DRAFT
		0	post#10	Title of Post 10	PUBLISHED

6. Get No. of Published Posts for a user

Primary key		Attributes		
Partition key: userId	Sort key: sk			
user_1	post#1	Title	Status	NoOfLikes
		Title of Post 1	PUBLISHED	10
	post#2	Title	Status	
		Title of Post 2	DRAFT	
	post#3	Title	Status	NoOfLikes
		Title of Post 3	PUBLISHED	7
	post#4	Title	Status	NoOfLikes
		Title of Post 4	PUBLISHED	12
	profile	Name	NoOfPublishedPosts	
		Rose Luettgen	3	
user_2	post#10	Title	Status	NoOfLikes
		Title of Post 10	PUBLISHED	0
	post#11	Title	Status	
		Title of Post 11	DRAFT	
	profile	Name	NoOfPublishedPosts	
		Ida Hegmann Sr.	1	

Transactions API

- Synchronous PUT, UPDATE, DELETE
- Atomic
- Automated Rollback





7. Get All Notifications for a user

Primary key		Attributes	
Partition key: userId	Sort key: sk		
user_1	notification#	notification#1	NotificationMessage
			Notification 1
		notification#2	NotificationMessage
			Notification 2
	post#1	Title	Status
		Title of Post 1	PUBLISHED
	post#2	Title	Status
		Title of Post 2	DRAFT
	post#3	Title	Status
		Title of Post 3	PUBLISHED
	post#4	Title	Status
		Title of Post 4	PUBLISHED
	profile	Name	NoOfPublishedPosts
		Rose Luettgen	3

userId,
BEGINS_WITH :
 "notification#"

8. Get Unread Notifications for a user

Index - {userId, NotificationsRead}

Indexes both unread and read

Primary key		Attributes	
Partition key: userId	Sort key: NotificationsRead		
user_1	false	sk	NotificationMessage
		notification#4	Notification 4
	true	sk	NotificationMessage
		notification#1	Notification 1
	true	sk	NotificationMessage
		notification#2	Notification 2
	true	sk	NotificationMessage
		notification#3	Notification 3

Delete NotificationsRead value for read notification

Primary key		Attributes		
Partition key: userId	Sort key: sk			
user_1	notification#1	NotificationMessage		
		Notification 1		
	notification#2	NotificationMessage		
		Notification 2		
	notification#3	NotificationMessage		
		Notification 3		
	notification#4	NotificationMessage	NotificationsRead	
		Notification 4	false	
	post#1	Title	Status	NoOfLikes
		Title of Post 1	PUBLISHED	10

Index - {userId, NotificationsRead}

Primary key		Attributes	
Partition key: userId	Sort key: NotificationsRead	sk	NotificationMessage
user_1	false	notification#4	Notification 4

9. Get No. of Unread Notifications for a user

Count

Single Table Design

userBlogApp	GSI: NoOfLikesIndex	GSI: UnreadNotificationsIndex	
Primary key		Attributes	
Partition key: userId	Sort key: sk		
user_1	notification#1	NotificationMessage	
		Notification 1	
	notification#2	NotificationMessage	
		Notification 2	
	notification#3	NotificationMessage	
		Notification 3	
	notification#4	NotificationMessage	NotificationsRead
		Notification 4	false
	post#1	Title	Status
		Title of Post 1	PUBLISHED
	post#2	Title	Status
		Title of Post 2	DRAFT
	post#3	Title	Status
		Title of Post 3	PUBLISHED
	post#4	Title	Status
		Title of Post 4	PUBLISHED
	profile	Name	NoOfPublishedPosts
		Rose Luetgen	3

Best Practices

- Identify your application's access patterns
- Understand the single-table design
- Avoid hot keys
- Avoid full-table scans
- Keep the sizes of indexes as small as possible
- Keep the number of indexes to a minimum

Data that is accessed together should be stored together.

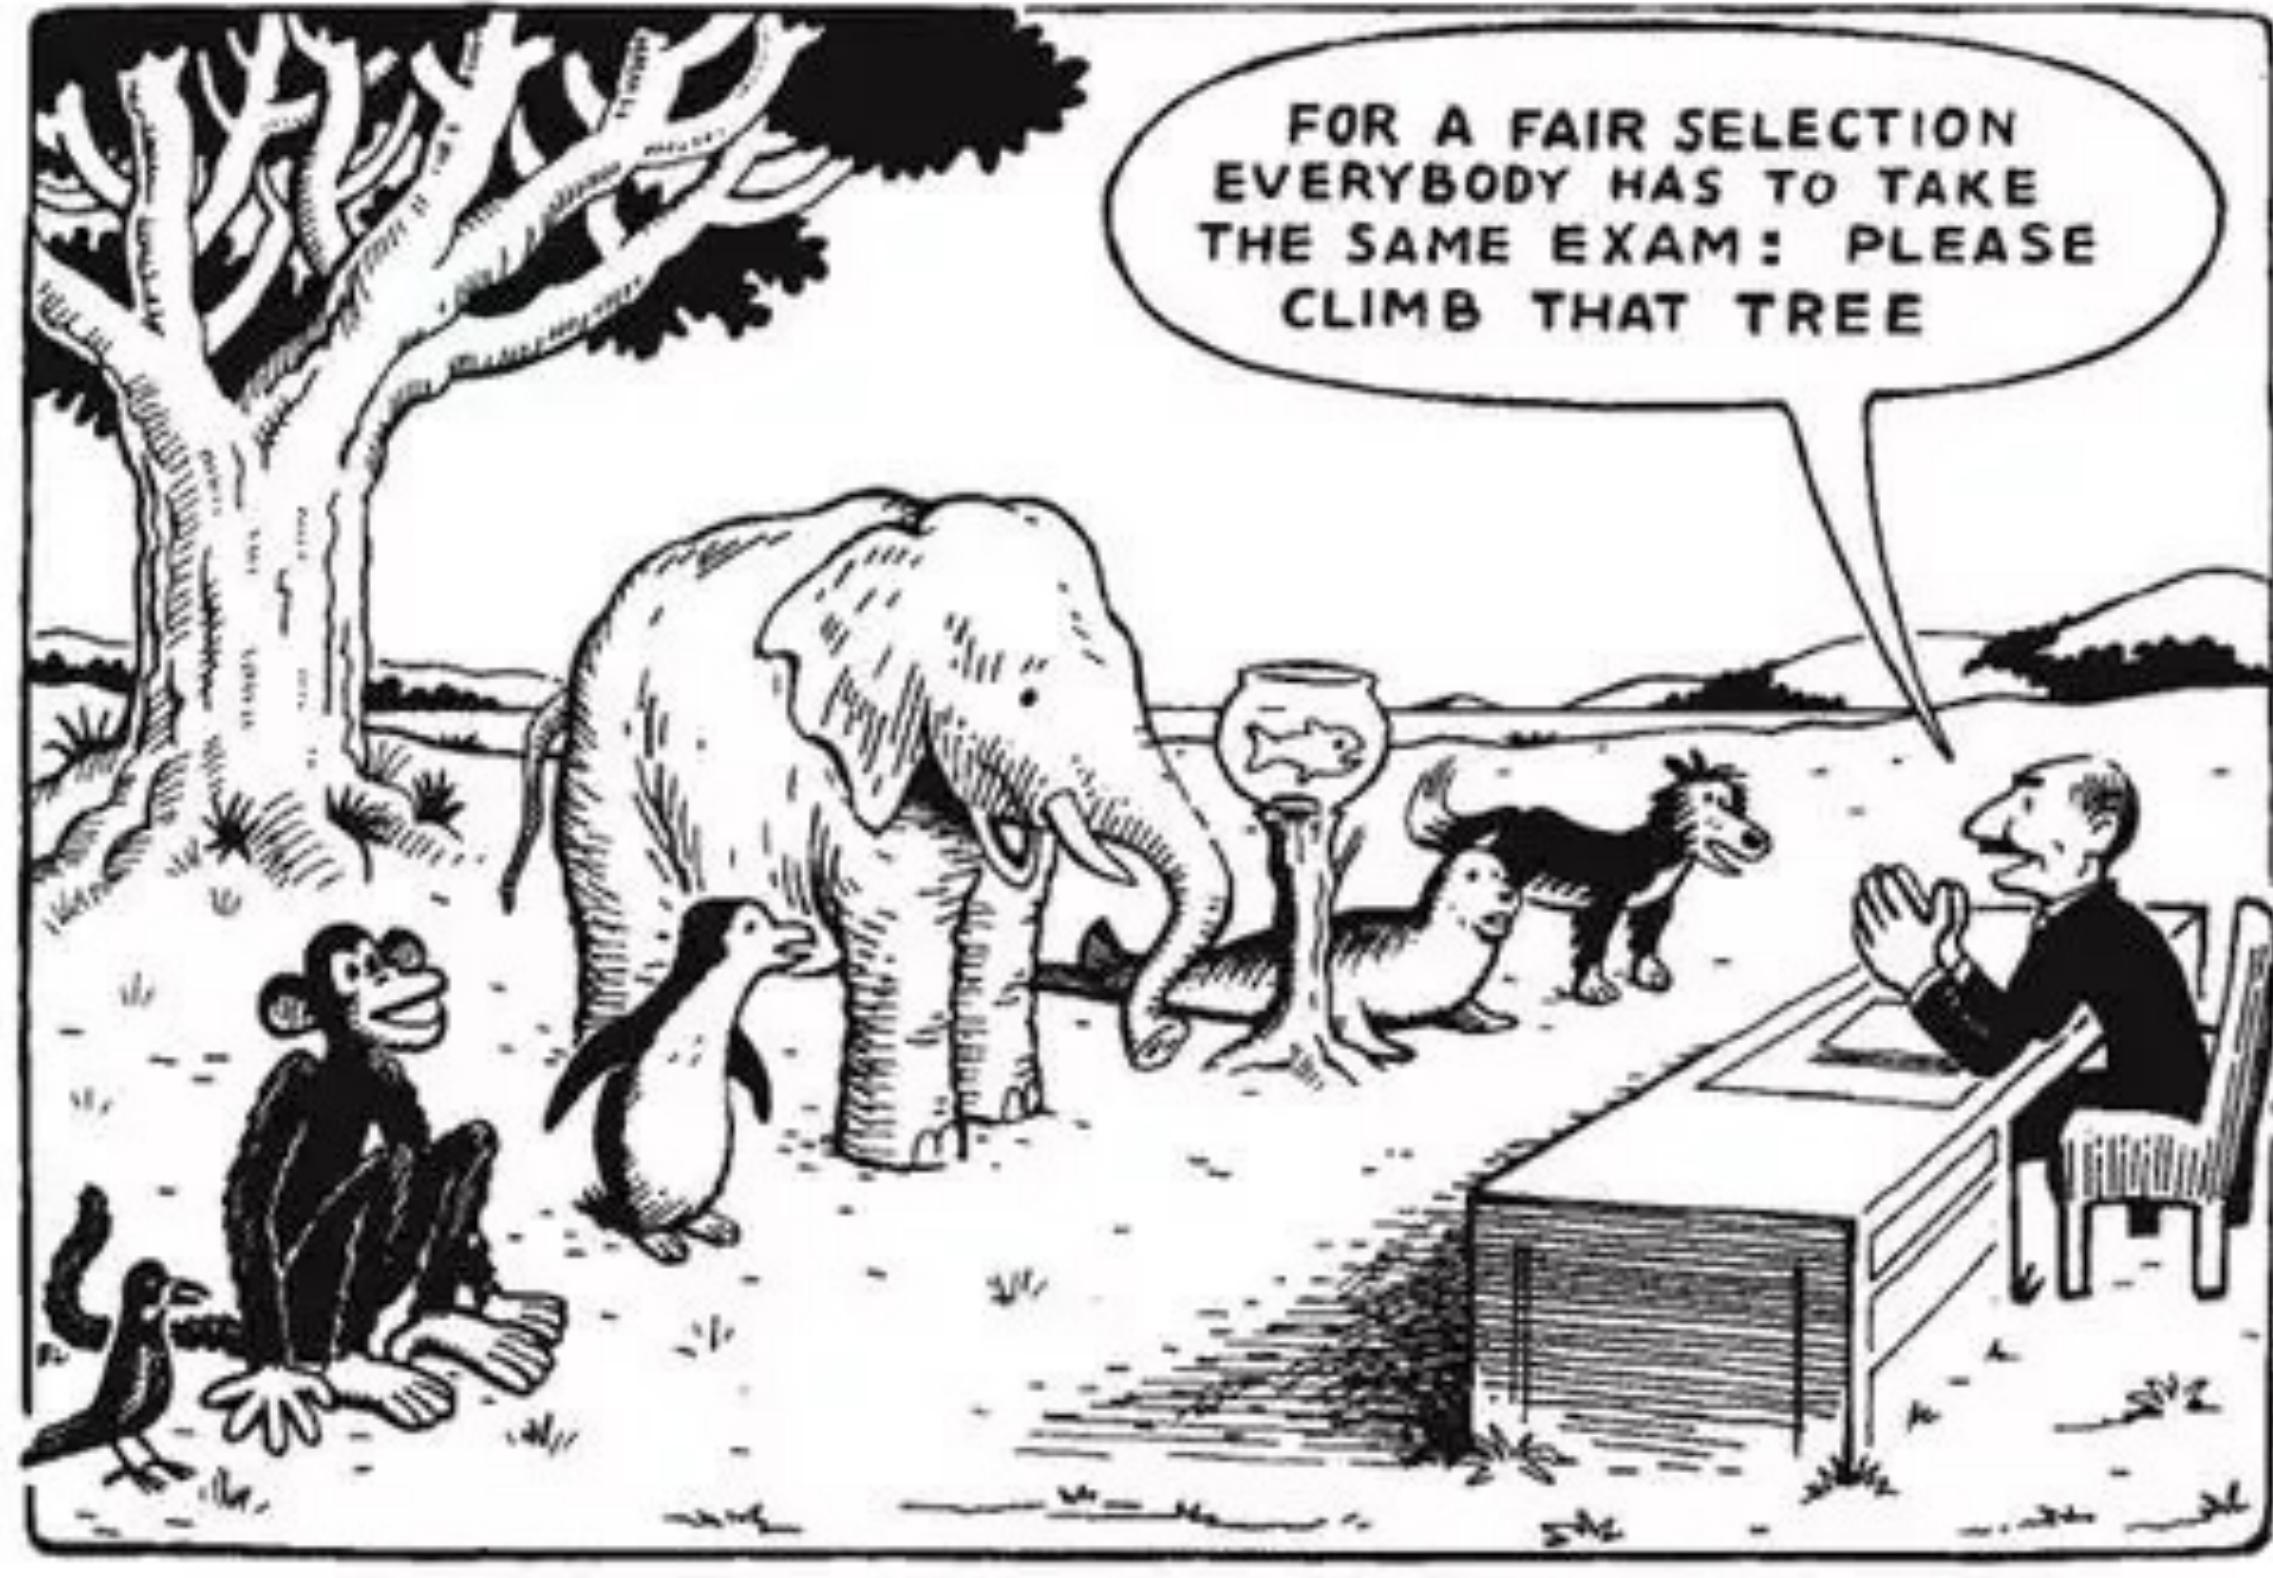






Dem

Content





Thank you! Happy No-SQLing!

Bhuvana

Technical Architect, Accenture



Bhuvana R

AWS Community Builder | 2x AWS Certified
| Architect | NodeJS





Gold Partner

PRESIDIO®

Silver Partner

ManageEngine
Site24x7
Crafted at Zoho Corp.

SAGENT

Securra Health

kijn
Where we belong

Syncfusion®

PRO